

cafe_feedback_exploration

November 19, 2025

1 Cafe Feedback Exploration & Walkthrough

This notebook captures the thinking, data processing steps, and exploratory visuals that underpin the Streamlit dashboard.

1.1 Walkthrough & Process

- **Goal:** surface hard (ratings, spend) and soft (comments) signals for cafe locations, and deliver actionable insights.
- **Steps taken:** inspect raw schema -> normalise Location strings -> parse/clean ratings and currency -> coerce timestamps -> derive calendar fields -> profile distributions for each column -> build visuals.
- **Why this matters:** the raw file contains stray spaces in locations and timestamp-like strings in Transaction Value; without cleaning, filters and spend averages become misleading. Keeping cleaning documented here keeps the dashboard uncluttered.
- **How to refresh:** rerun this notebook on new drops; cleaned outputs and charts will stay in sync with the app.

1.2 Quick explanation of data processing

- Drop empty Unnamed columns; coerce Rating to numeric.
- Strip/collapse whitespace in Location to merge duplicates (e.g., Albany -> Albany).
- Parse Transaction Value as currency, strip \$/commas, and discard impossible spends (<=0 or >) to remove timestamp artefacts.
- Parse Transaction Date and Time as day-first datetime; derive Date and DayName for trend views.
- Persist these steps so dashboard filters and aggregates stay realistic.

```
[48]: import re
from pathlib import Path

import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

sns.set_theme()

DATA_FILE = Path("Sample data - Cafe - Sample data 2400 records.csv")
raw_df = pd.read_csv(DATA_FILE)
```

```

raw_df.head()

[48]:          Location  Rating  \
0           Domain      5
1        Timaru      5
2   Upper Hutt      5
3 Wellington central      4
4     Henderson      5

                                         Comment  \
0  Awesome, friendly staff, they all are smiling ...
1                           The service
2  The staff are fabulous and neither is too much...
3            More dairy + gluten free food options.
4  Friendly welcoming staff always smiling

  Transaction Date and Time Transaction Value feedback_id Unnamed: 6  \
0  19/10/2024 6:03:00 AM       $19.50    4974583      NaN
1  19/10/2024 6:09:00 AM       $19.17    4974584      NaN
2  19/10/2024 6:13:00 AM       $10.60    4974586      NaN
3  19/10/2024 6:21:00 AM        $4.00    4974587      NaN
4  19/10/2024 6:43:00 AM       $5.50    4974591      NaN

  Unnamed: 7
0      NaN
1      NaN
2      NaN
3      NaN
4      NaN

```

1.3 Raw exploration before cleaning

- Visualise raw values first to spot anomalies that drive cleaning decisions.

```

[49]: # Raw distributions prior to cleaning
raw_location_counts = raw_df["Location"].value_counts()
print("Locations")
display(raw_location_counts)

# Duplicate locations due to extra spaces
normalized_locations = raw_df.assign(
    NormalizedLocation=raw_df["Location"].astype(str).str.strip())
)
dup_location_table = (
    normalized_locations.groupby("NormalizedLocation", dropna=False)
    .agg(
        sample_raw=("Location", "first"),

```

```

        occurrences=("Location", "size"),
        unique_raw_variants=("Location", "nunique"),
    )
    .reset_index()
    .rename(
        columns={
            "NormalizedLocation": "Location (duplicate)",
            "sample_raw": "Raw location sample",
            "occurrences": "Occurrences",
            "unique_raw_variants": "Unique raw variants",
        }
    )
)
dup_location_table = (
    dup_location_table[dup_location_table["Unique raw variants"] > 1]
    .sort_values("Occurrences", ascending=False)
    .drop(columns="Unique raw variants")
)
print("\nDuplicated location table:")
display(dup_location_table)

# Transaction Value: non-currency patterns
raw_txn_strings = raw_df["Transaction Value"]
print("\nExamples of transaction values containing non-currency patterns:")
display(raw_txn_strings[raw_txn_strings.str.contains("/", na=True)])

# Transaction Value: numeric outliers
raw_txn_numeric = pd.to_numeric(
    raw_txn_strings.str.replace("[^0-9.]", "", regex=True), errors="coerce"
)
anomaly_upper_threshold = 500
anomaly_lower_threshold = 0
anomalies = raw_txn_numeric[
    (raw_txn_numeric > anomaly_upper_threshold) | (raw_txn_numeric <= anomaly_lower_threshold)
].sort_values(ascending=False)
print(f"{len(anomalies)} raw transaction values are outside the expected range\n"
      f"[{0}, {anomaly_upper_threshold}]")
display(anomalies)

# Ratings: value counts and out-of-range/missing
rating_raw = raw_df['Rating']
print('Rating value counts (raw):')
display(rating_raw.value_counts(dropna=False))
rating_numeric = pd.to_numeric(rating_raw, errors='coerce')

```

```

rating_anomalies = rating_raw[(rating_numeric.isna()) | (rating_numeric < 1) | (rating_numeric > 5)]
print(f"{len(rating_anomalies)} ratings are missing or outside 1-5")
display(rating_anomalies)

# Transaction Date and Time: non-parsable examples
raw_dates = raw_df["Transaction Date and Time"]
standard_format = "%d/%m/%Y %I:%M:%S %p"
parsed_dates = pd.to_datetime(raw_dates, format=standard_format, errors="coerce")
bad_dates = raw_dates[parsed_dates.isna()]
print(f"{len(bad_dates)} transaction datetime values could not be parsed")
display(bad_dates)

# feedback_id: blanks, non-numeric, duplicates
feedback_raw = raw_df['feedback_id'].astype(str).str.strip()
bad_feedback = feedback_raw[(feedback_raw == '') | (~feedback_raw.str.match(r'^\d+$'))]
print(f"{len(bad_feedback)} feedback_id values are non-numeric")
display(bad_feedback)
dup_feedback = feedback_raw[feedback_raw.duplicated(keep=False)]
print(f"{len(dup_feedback)} feedback_id values are duplicates")
display(dup_feedback)

# Blank unnamed columns
cols_to_display = raw_df.columns[6:9]
print(f"Inspecting columns (6th & 7th): {cols_to_display.tolist()}")
preview_cols = (
    raw_df[cols_to_display]
    .astype("object")
    .sort_values(by=cols_to_display.tolist(), ascending=False)
    .head(10)
)
display(preview_cols)

```

Locations

Location	
Wanganui	54
Rangiora	51
Palmerston Nth	48
Palmerston Nth	47
Kapiti	47
	..
Whakatane	15
Henderson	12
Pukekohe	11
Lincoln Road	7

```
Queen Street      2
Name: count, Length: 75, dtype: int64
```

Duplicated location table:

	Location (duplicate)	Raw location sample	Occurrences
40	Palmerston Nth	Palmerston Nth	95
22	Invercargill	Invercargill	77
34	Nelson	Nelson	77
5	Botany	Botany	63
51	Rotorua	Rotorua	60
0	Albany	Albany	55
44	Pukekohe	Pukekohe	49
20	Henderson	Henderson	38

Examples of transaction values containing non-currency patterns:

```
509    20/10/2020 8:24:00 AM
Name: Transaction Value, dtype: object

92 raw transaction values are outside the expected range [0, 500]

509    2.010202e+12
2362   2.410201e+11
1986   2.310201e+11
1556   2.210201e+11
1249   2.210201e+11

...
2269   0.000000e+00
2293   0.000000e+00
2340   0.000000e+00
2367   0.000000e+00
2393   0.000000e+00

Name: Transaction Value, Length: 92, dtype: float64
```

Rating value counts (raw):

```
Rating
5     1980
4     335
3     68
2     16
1     8

Name: count, dtype: int64
```

0 ratings are missing or outside 1-5

Series([], Name: Rating, dtype: int64)

5 transaction datetime values could not be parsed

```

509                      WEB
1249      kim.wise@auckland.ac.nz
1556      robsued47@gmail.com
1986      tkmolloy@gmail.com
2362      lynn.m@xtra.co.nz
Name: Transaction Date and Time, dtype: object

5 feedback_id values are non-numeric

509      16.1
1249      WJZPJCF
1556      NFTUTGH
1986      MTKPDCF
2362      MFFPYCF
Name: feedback_id, dtype: object

0 feedback_id values are duplicates

Series([], Name: feedback_id, dtype: object)

Inspecting columns (6th & 7th): ['Unnamed: 6', 'Unnamed: 7']

    Unnamed: 6  Unnamed: 7
509      DEPLHCF  4975742.0
2362      4979455      NaN
1986      4978775      NaN
1556      4977933      NaN
1249      4977334      NaN
0          NaN      NaN
1          NaN      NaN
2          NaN      NaN
3          NaN      NaN
4          NaN      NaN

```

Raw-data findings: - Locations appear multiple times due to extra space (e.g., Albany vs Albany). - Transaction Value includes timestamp-like strings (e.g., 20/10/2020 8:24:00 AM) and extreme magnitudes (e.g., 241020120707.0) that would inflate spend if unfiltered. - Transaction Date and Time has entries that are not datetime. - feedback_id includes non-numeric strings that should be cleaned. - 2 unnamed columns are almost entirely blank and can be removed. - These observed issues motivate the cleaning rules applied below.

```
[50]: # Cleaning helpers
def parse_transaction(val):
    if pd.isna(val):
        return None
    s = str(val).strip()
    m = re.search(r"-?\$?\s*([0-9]{1,3}(?:,[0-9]{3})*|[0-9]+)(?:\.[0-9]{1,2})?"
    ↵", s)
    if not m:
        return None
    try:
```

```

        v = float(re.sub(r"[,$]", "", m.group()))
    except ValueError:
        return None
    if v <= 0 or v > 500:
        return None
    return v

df = raw_df.copy()
df = df.loc[:, ~df.columns.str.contains("^\u0331Unnamed")]
df["Location"] = df["Location"].astype(str).str.strip().str.replace(r"\s+", " ", regex=True)
df["Rating"] = pd.to_numeric(df["Rating"], errors="coerce")
df["Transaction Value"] = df["Transaction Value"].apply(parse_transaction)
df["Transaction Date and Time"] = pd.to_datetime(df["Transaction Date and Time"], dayfirst=True, errors="coerce")
df["Date"] = df["Transaction Date and Time"].dt.date
df["DayName"] = df["Transaction Date and Time"].dt.day_name()
df.dropna(subset=["Rating", "Transaction Value"], inplace=True)
df.head()

```

[50]:

	Location	Rating	Comment	Date
0	Domain	5	Awesome, friendly staff, they all are smiling ...	2024-10-19
1	Timaru	5	The service	2024-10-19
2	Upper Hutt	5	The staff are fabulous and neither is too much...	2024-10-19
3	Wellington central	4	More dairy + gluten free food options.	2024-10-19
4	Henderson	5	Friendly welcoming staff always smiling	2024-10-19

	Transaction Date and Time	Transaction Value	feedback_id	Date
0	2024-10-19 06:03:00	19.50	4974583	2024-10-19
1	2024-10-19 06:09:00	19.17	4974584	2024-10-19
2	2024-10-19 06:13:00	10.60	4974586	2024-10-19
3	2024-10-19 06:21:00	4.00	4974587	2024-10-19
4	2024-10-19 06:43:00	5.50	4974591	2024-10-19

	DayName
0	Saturday
1	Saturday
2	Saturday
3	Saturday
4	Saturday

1.4 Categorical distributions (all categories printed)

For each categorical-like column, the full value counts are displayed. Plots are shown when the category count is manageable.

```
[51]: categorical_cols = []
numeric_cols = []

for col in df.columns:
    if pd.api.types.is_numeric_dtype(df[col]):
        numeric_cols.append(col)
    elif pd.api.types.is_datetime64_any_dtype(df[col]):
        categorical_cols.append(col)
    else:
        categorical_cols.append(col)

for col in categorical_cols:
    print(f"\n==== {col} (categorical-like) ====")
    vc = df[col].fillna("<missing>").value_counts()
    display(vc)
    if 1 < len(vc) <= 30:
        vc.sort_values().plot(kind="barh", title=f"{col} frequency", □
        ↵ figsize=(8, max(4, len(vc) * 0.25)))
        plt.xlabel("Count")
        plt.tight_layout()
        plt.show()
```

==== Location (categorical-like) ====

```
Location
Palmerston Nth    92
Invercargill     76
Nelson            75
Botany             62
Rotorua            59
                    ..
Williams Drive    16
Rosebank Road     15
Whakatane          14
Lincoln Road       7
Queen Street        1
Name: count, Length: 67, dtype: int64
```

==== Comment (categorical-like) ====

```
Comment
```

Great service		□
↳		□
↳	14	□
Friendly staff		□
↳		□
↳	13	□
Good service		□
↳		□
↳	9	□
Service		□
↳		□
↳	8	□
Friendly service		□
↳		□
↳	8	□
↳		□
↳	..	□
The Muffin was very tasty.		□
↳		□
↳	1	□
love the friendly staff at the Nelson shop, they're really awesome!		□
↳		□
↳	1	□
Engagement with the wait staff was warmer and friendly.		□
↳		□
↳	1	□
Rob and the team are so lovely they welcomed me with a big smile and by name . I go there all the time and they know what coffee I like and it have ready for me by the time I get up to the counter . 1		□
Always clean with good selection of food. Fantastic coffee with friendly personal service.		□
↳	1	□
Name: count, Length: 2185, dtype: int64		

==== Transaction Date and Time (categorical-like) ====

Transaction Date and Time

<missing>	199	
2024-10-24 09:28:00	7	
2024-10-19 08:50:00	6	
2024-10-21 09:27:00	6	
2024-10-19 09:44:00	6	
...		
2024-10-24 11:13:00	1	
2024-10-24 09:19:00	1	
2024-10-23 08:57:00	1	

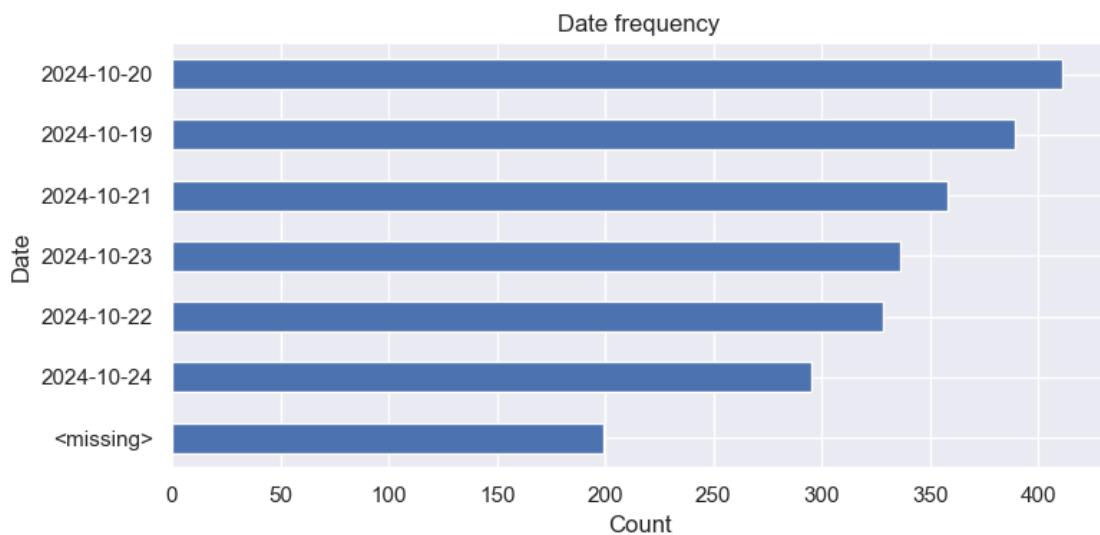
```
2024-10-22 10:22:00      1  
2024-10-24 10:14:00      1  
Name: count, Length: 1260, dtype: int64
```

===== feedback_id (categorical-like) =====

```
feedback_id  
4979626      1  
4974583      1  
4974584      1  
4974586      1  
4974587      1  
..  
4974621      1  
4974622      1  
4974623      1  
4974625      1  
4974626      1  
Name: count, Length: 2316, dtype: int64
```

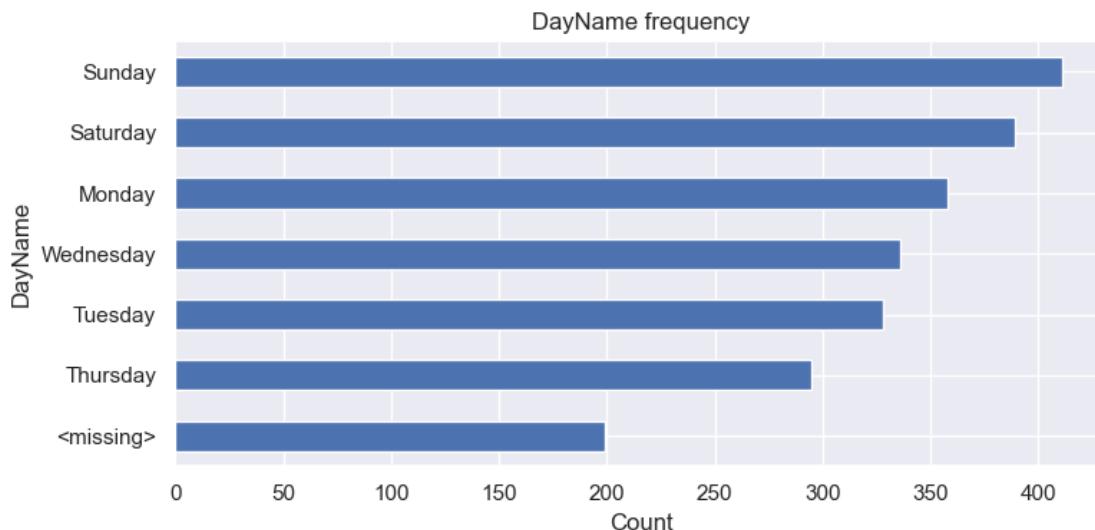
===== Date (categorical-like) =====

```
Date  
2024-10-20    411  
2024-10-19    389  
2024-10-21    358  
2024-10-23    336  
2024-10-22    328  
2024-10-24    295  
<missing>    199  
Name: count, dtype: int64
```



```
==== DayName (categorical-like) ====
```

```
DayName
Sunday      411
Saturday    389
Monday      358
Wednesday   336
Tuesday     328
Thursday    295
<missing>   199
Name: count, dtype: int64
```



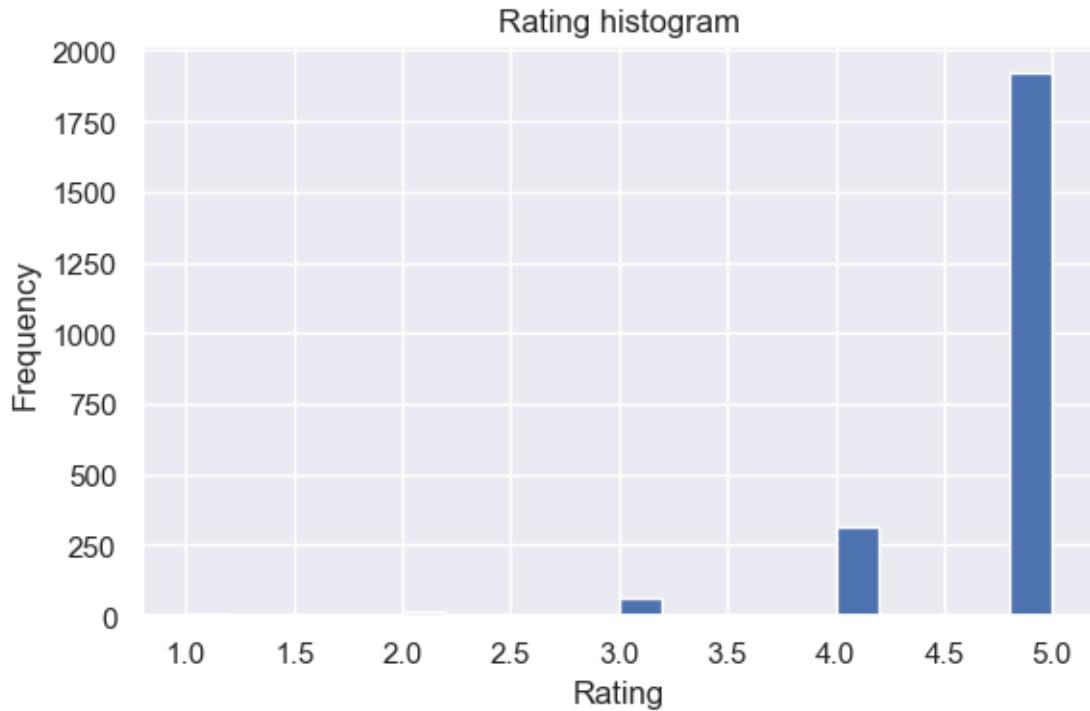
1.5 Numeric distributions

Descriptive stats and histograms for each numeric-like column.

```
[52]: for col in numeric_cols:
    print(f"\n==== {col} (numeric-like) ====")
    series = df[col].dropna()
    display(series.describe())
    series.plot(kind="hist", bins=20, title=f"{col} histogram", figsize=(6,4))
    plt.xlabel(col)
    plt.tight_layout()
    plt.show()
```

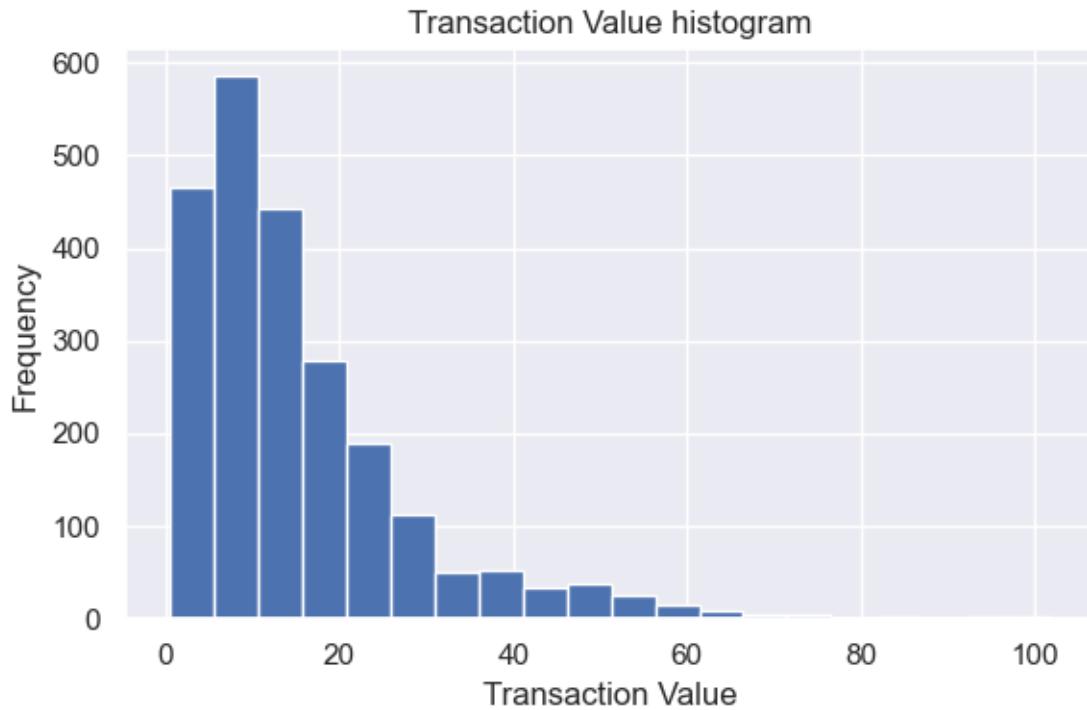
```
==== Rating (numeric-like) ====
```

```
count    2316.000000
mean     4.779793
std      0.543445
min     1.000000
25%     5.000000
50%     5.000000
75%     5.000000
max     5.000000
Name: Rating, dtype: float64
```



```
===== Transaction Value (numeric-like) =====
```

```
count    2316.000000
mean     15.967414
std      12.962577
min     0.500000
25%     6.800000
50%     11.500000
75%     20.000000
max     102.000000
Name: Transaction Value, dtype: float64
```



1.6 Targeted visuals

- Rating distribution and average transaction by rating
- Daily trends of rating and spend (to mirror the dashboard)
- Top comment keywords

```
[53]: rating_counts = df["Rating"].value_counts().sort_index()
rating_counts.plot(kind="bar", title="Rating distribution", ylabel="Count", xlabel="Rating")
plt.tight_layout()
plt.show()

avg_txn_by_rating = df.groupby("Rating")["Transaction Value"].mean()
avg_txn_by_rating.plot(kind="bar", title="Avg transaction by rating", ylabel="$", xlabel="Rating")
plt.tight_layout()
plt.show()

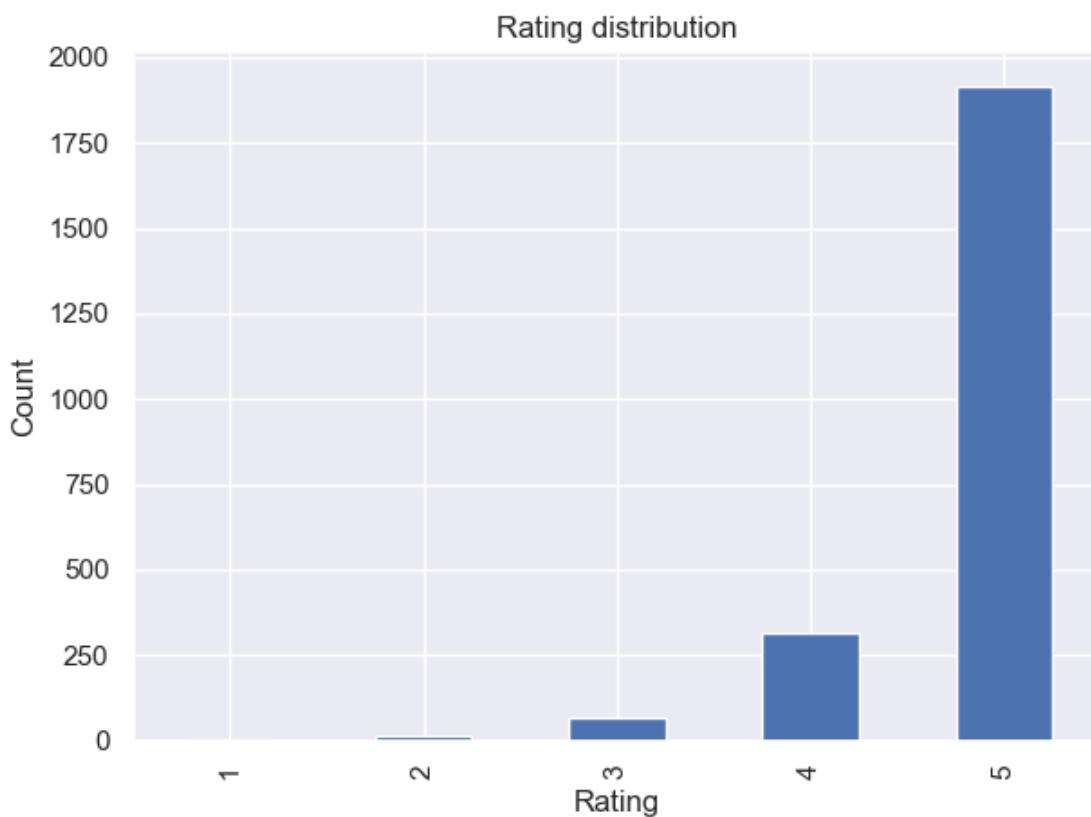
daily = (
    df.groupby("Date")
    .agg(avg_rating=("Rating", "mean"), avg_txn=("Transaction Value", "mean"))
    .sort_index()
)
```

```

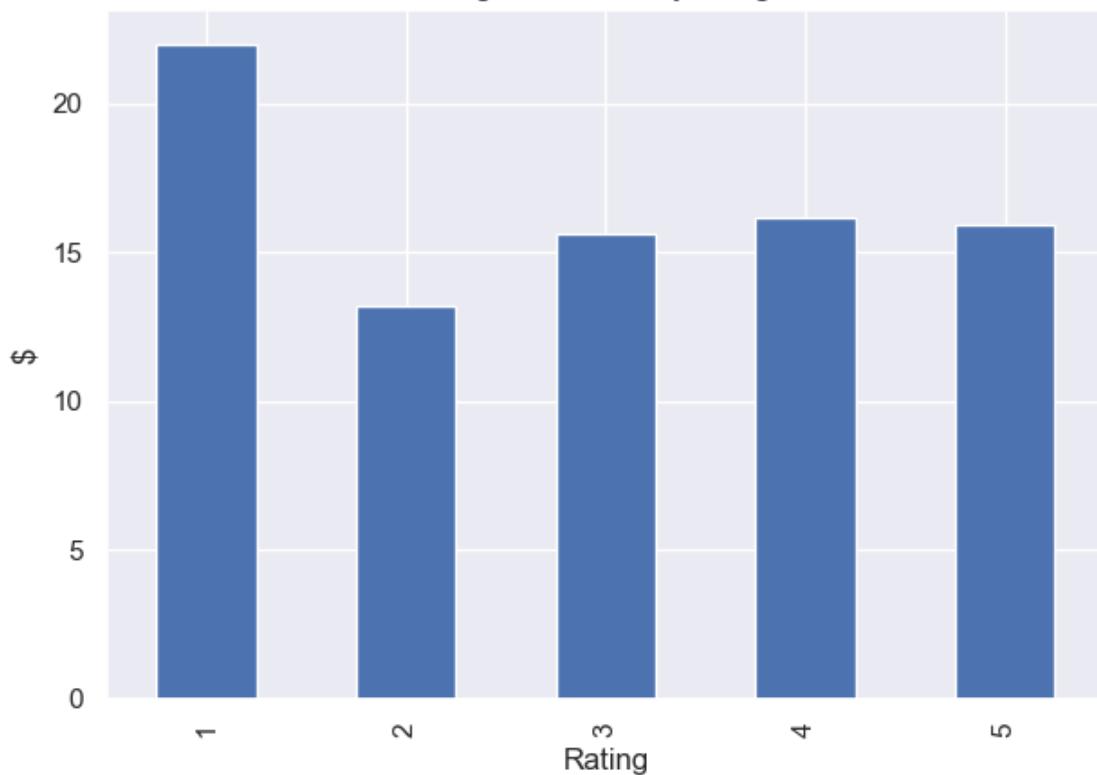
daily.plot(y=["avg_rating", "avg_txn"], title="Daily trends: rating and spend",
           marker="o")
plt.tight_layout()
plt.show()

stop_words = [
    "the", "and", "to", "of", "a", "in", "for", "with", "is", "it", "on", "my", "our", "at", "are", "was", "be",
    "word_counts = {}
for c in df["Comment"].dropna().astype(str):
    for w in re.findall(r"[a-zA-Z]+", c.lower()):
        if w not in stop_words and len(w) > 1:
            word_counts[w] = word_counts.get(w, 0) + 1
top_words = sorted(word_counts.items(), key=lambda kv: kv[1], reverse=True)[:12]
pd.DataFrame(top_words, columns=["word", "count"]).set_index("word").
    plot(kind="bar", title="Top words in comments")
plt.tight_layout()
plt.show()

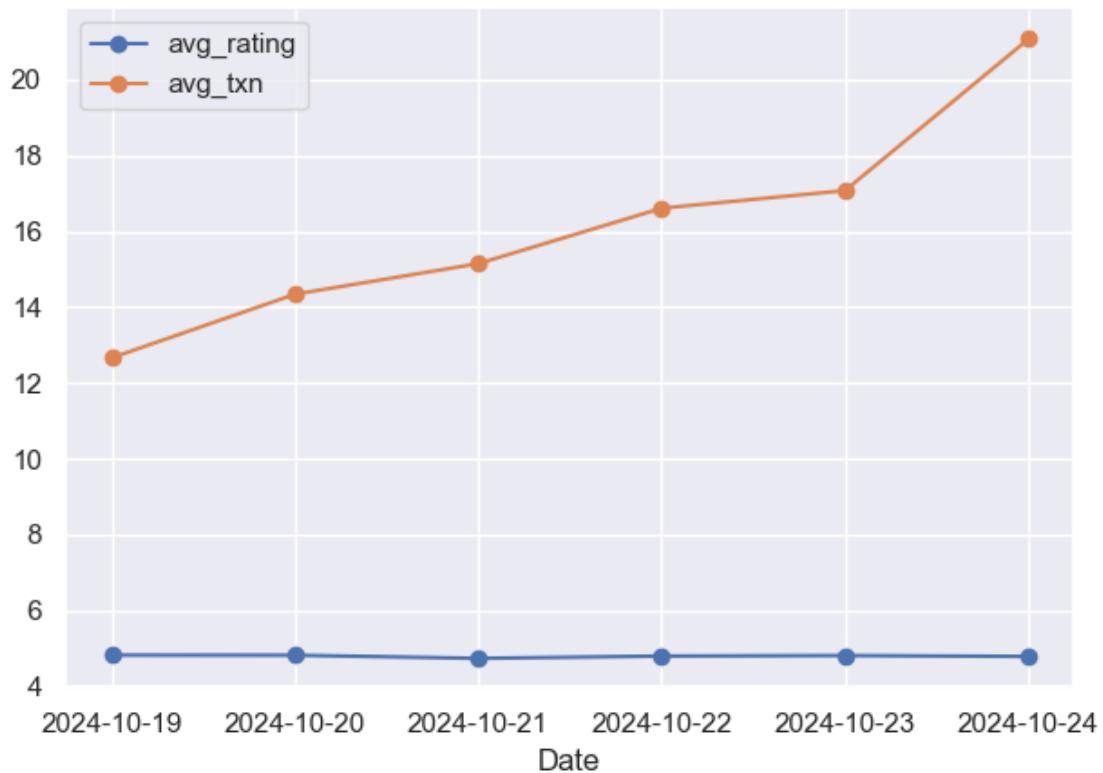
```

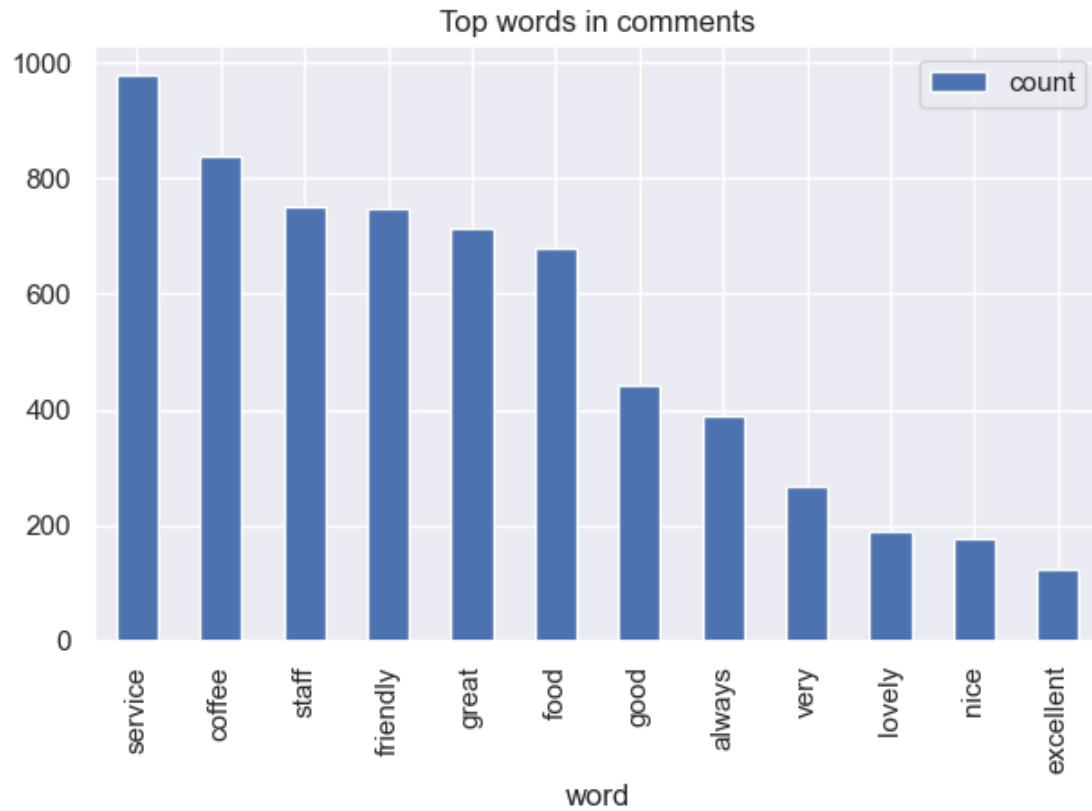


Avg transaction by rating



Daily trends: rating and spend





1.7 Key summaries

- Location labels need normalisation; otherwise filters show duplicates.
- Cleaning transaction values removes timestamp artefacts and keeps spend realistic (~\$15 median).
- Ratings skew high while low ratings are scarce; focus on consistency and speed where comments flag issues.
- Daily trend views help spot emerging changes faster than monthly rolls.